

1 Notations

This write-up presents my PVI and PNVI memory safety policies. In the interest of being able to slot it into the eventual paper and thesis, I'm assuming that we have given the standard explanations of what tags and control points are. I do need to get into the semantics and notations a bit, however.

Values are ranged over by v , variable identifiers by x , and function identifiers by f . Tags use a number of metavariables: vt , pt , lt , nt , and \mathcal{P} . An *atom* is a pair of a value and a tag, $v@vt$; the @ symbol should be read as a pair in general, and is used when the second object in the pair is a tag. Expressions are ranged over by e , statements by s , and continuations by k . Many continuations contain a context ctx , which is a function from expressions to expressions.

Global environments, ranged over by ge , map identifiers to either function or global variable definitions (including the variable's location in memory. Local environments, ranged over by le , map identifiers to atoms. And the *label environment* Memories m map integers to triples: a value, a "value tag" vt , and a list of "location tags" \overline{lt} .

A memory is an array of bytes, and a load or store will access some number of bytes. I write $m[l]_s = v@vt@\overline{lt}$ to denote loading s bytes, starting at location l , and interpreting them as a value v , a value tag vt , and a list of s location tags. Likewise, $m[l \mapsto v@vt@\overline{lt}]_s$ denotes storing that many bytes. vt is tied to a full value, which may consist of multiple bytes, while each tag in \overline{lt} is tied to an individual byte. When writing multiple contiguous values, I will write a range of locations. So in the case of an array of 10 integers, s would be 4, and $m[l \dots l + 10 \mapsto v@vt@\overline{lt}]_4$ would write $v@vt$ to ten words starting at l , with the four bytes of each word tagged with \overline{lt} 's four tags. This guarantees that even misaligned loads and stores always have a valid location tag to check (possibly multiple, mismatched location tags, in which case the policy can failstop if needed.)

State can be of several kinds, denoted by their script prefix: a *general state* $\mathcal{S}(\dots)$, an *expression state* $\mathcal{E}(\dots)$, a *call state* $\mathcal{C}(\dots)$, or a *return state* $\mathcal{R}(\dots)$. Finally, the special state *failstop* (\mathcal{F}) represents a tag failure.

$$\begin{aligned} S ::= & \mathcal{S}(m, ge, le \mid s;; k@\mathcal{P}) \\ & \mid \mathcal{E}(m, ge, le \mid k[e : ty]@\mathcal{P}) \\ & \mid \mathcal{C}(m, ge, le \mid k[f(\overline{v@vt})]@\mathcal{P}) \\ & \mid \mathcal{R}(m, ge, le \mid k[v@vt]@\mathcal{P}) \\ & \mid \mathcal{F} \end{aligned}$$

Control Points with Side-effects and Optional Arguments Most control points can be mapped cleanly onto one or more instructions in a compiled program. For example, the **BinopT** control point takes as input the tags on the parameters of an operation (as well as the PC tag) and yields a tag for the result, so the target-level rule, which does the same, can be identical. Other control points may correspond to multiple target-level instructions, requiring a more complicated mapping. I will not call these out unless they are particularly noteworthy. From a performance standpoint, the most problematic situation is when a Tagged-C control point requires a tag from a location that is not read under a normal compilation scheme, which must update tags in locations that are not written, or in which the source construct does not have corresponding instructions in the target.

These situations require the compiler to add instructions for the primary purpose of manipulating tags. If the tag rules that instantiate those control points do not make use of them, these

instructions are needless overhead. In these cases, the control points will take optional parameters or return optional results, and I will explain how the rule should be implemented in the target if the options are used. If compiling with a known policy that does not make use of the options, it will be sound to eliminate the extra instructions. If *all* of the control point’s outputs are optional and unused, the control point need not be compiled at all. In this document, optional inputs and outputs will be marked with boxes.

Name Tags When we want to define a per-program policy, we need to be able to attach tags to the program’s functions, globals, and so on. We do this by automatically embedding their identifiers in tags, which are available to all policies. These are called *name tags* and are ranged over by *nt*. The complete list of such tags is as follows:


- Function identifiers
- Function arguments, written `f.x`
- Global variables

2 Memory safety basics

Memory safety policies operate on under a “lock and key” model, in which objects in memory are tagged with a unique identifier (the “lock”) and may only be accessed via a pointer tagged with the same identifier (the “key.”) For a simple example, consider the following code:

```
void main() {  
    int a[10];  
    int b[10];  
    a[10] = 42;  
}
```

In a typical stack allocator—such as the one used by my interpreter—`a` and `b` will be allocated next to one another on the stack, like this:



To prevent the expression `a[10] = 42` from overwriting `b[0]`, we give `a` and `b` unique *color tags* when they are allocated. In this case, we’ll tag `a` with *dyn 0*, indicating that it’s the first dynamically allocated object, and `b` with *dyn 1*. Then, when we evaluate the left-hand expression `a` into its memory location *l*, we tag *l* with *dyn 0*. When we take the offset *l* + 10, we keep that tag. And when we perform the assignment, we check that the location tag at *l* matches. It doesn’t, so we failstop.

The same principle applies for this code:

```

void main() {
    int* a = malloc(10 * sizeof(int));
    int* b = malloc(10 * sizeof(int));
    *(a + (b - a)) = 42;
}

```

In this case, `a` and `b` could be allocated anywhere in the heap, and in Tagged C the expression `*(a + (b - a)) = 42` will always write to `*b`. While this might be intentional on the part of the programmer, it is also undefined behavior in the C standard, and in some (but not all; see below) formal C semantics. Likewise, if `a` and `b` are next to each other or in some other predictable arrangement, arithmetic like our first example can apply. The memory safety policy works just the same in this scenario, with the tags being attached by the call to `malloc`, once again using the *dyn* label in a global count of allocated blocks. Meanwhile, values that are not derived from valid pointers at all are tagged `X`, and can never be read or written through, to avoid pointer forging, like this:

```

void main() {
    int* a = malloc(10 * sizeof(int));
    // I happen to know that a will be at address 1000
    *1000 = 42;
}

```

Both stack and heap allocations use the *dyn* label and have a color that can grow arbitrarily high. This is because over a program’s execution, it might allocate an unbounded number of heap- or stack-allocated objects, and each needs a unique identifier. Existing work has shown that in practice, tag colors can be “garbage collected” and reused, but in Tagged C we assume them to be infinite and unique.

Lastly, we have global variables. While “global safety” is not as prominent a topic as heap or stack safety, overrunning a global buffer is still a problem. It is also easy to forge a pointer to a global, and when this happens it can undermine assumptions about the behavior of linked libraries whose globals are not exported. Globals do not need dynamic colors, but can use their identifiers as tags, of the form *glob id*.

3 Memory Safety: PVI and PNVI

Our policies aim to enforce two memory models in particular: *PVI* (provenance via integer) and *PNVI* (provenance not via integer) from Memarian et al. [?]. They propose PVI and PNVI as memory models that support common idioms that are undefined in the C standard, but are still restrictive enough as to support useful alias analysis for optimization. This application is orthogonal to security, and violations of either memory model are treated as undefined behavior, just as in the C standard. Our goal is to turn that UB into failstop behavior, so that undefined programs cannot accidentally undermine their own security.

Both memory models represent pointers as integers, just as Tagged C does, with additional provenance associated with each object. An integer cast to a pointer in PVI retains this provenance,

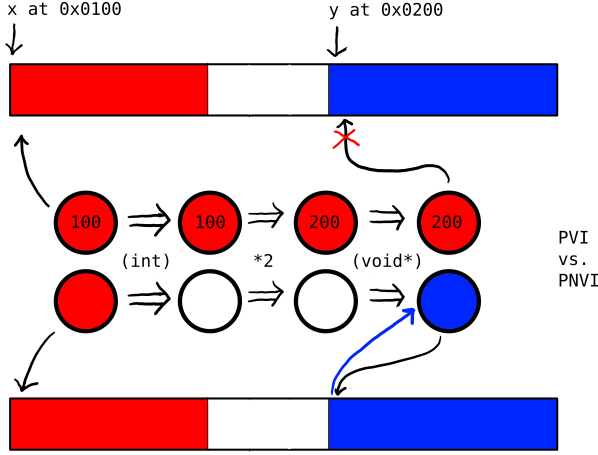


Figure 1: Integer-pointer casts in PVI and PNVI

enabling integer operations to be performed on it prior to it being cast back to a pointer. In PNVI, by contrast, an integer cast to a pointer gains the provenance of the object it points to when the cast occurs. While PNVI supports a wider range of programs, it is inconsistent with important assumptions of the C memory model, in ways that may have serious security consequences. The difference between PVI and PNVI is illustrated in Figure 1.

We will aim to prove that for any program, if it is run in both the PVI semantics and in Tagged C with our PVI policy, it either produces identical output, or it is both undefined in the PVI semantics and failstops in Tagged C. Likewise for PNVI, except that some UB in PNVI is non-deterministic, and we only require that it failstop in an execution that would *reach* the UB.

3.1 PVI Definitions

Here we give the relevant tag rules for the PVI policy, and describe the control points that they attach to. We will, for each rule, first give the control point(s) that use it, along with a brief explanation of what the surrounding semantics rules do, and then give the rule. For these policies, all control points appear in expression reduction steps. The machine state consists of the PC tag \mathcal{P} , a memory m , the global environment ge , and a local environment le . These are contextual semantics, so each expression is situated in some context ctx .

The core of the PVI policy is the *provenance color*, represented by a natural number.

$$\begin{array}{ll} T ::=_{dyn} glob\ id & id \in ident \\ & C \in \mathbb{N} \end{array}$$

Color generation New colors are generated when objects are allocated. When exactly that occurs depends on where the object lives. Global variables are a special case: they are allocated

during program initialization, before execution begins. As such they do not have a control point per se, but a rule that functions similarly, while being more expressive.

Given a list xs of variable identifiers id and types ty , a program's initial memory is defined by iteratively allocating each one in memory and updating the global environment with its base address, bound, type, and a static identity tag. Let $|ty|$ be a function from types to their sizes in bytes. The memory is initialized **undef@vt@ \overline{lt}** for some vt and \overline{lt} , unless given an initializer. Let m_0 and ge_0 be the initial (empty) memory and environment. The parameter b marks the start of the global region.

$$globals\ xs\ b = \begin{cases} (m_0, ge_0) & \text{if } xs = \varepsilon \\ (m[p \dots p + |ty| \mapsto \mathbf{undef@vt@}\overline{lt}]_{|ty|}, & \text{if } xs = (id, ty) :: xs' \\ ge[id \mapsto (p, p + |ty|, ty, pt)]) & \text{and } pt, vt, \overline{lt} \leftarrow \mathbf{GlobalT}(id, |ty|) \\ & \text{where } (m, ge) = globals\ xs' (b + |ty|) \end{cases}$$

GlobalT(id, s)

$$\begin{aligned} pt &\leftarrow glob\ id \\ vt &\leftarrow X \\ \overline{lt} &\leftarrow [glob\ id \mid 0 \leq i < s] \end{aligned}$$

Stack-allocated locals are allocated at the start of a function call. Like a global environment, a local environment maps identifiers to base, bound, type, and tag. The rule is almost identical to allocation of globals, except that the stack allocator, *stack_alloc* will be more complex in order to support deallocation (in practice, it uses a normal stack structure and allocates and deallocates by increasing and decreasing a “stack pointer”).

Since allocations occur at runtime, the value and location tags that initialize the allocated memory are optional. They would be realized by initializing the entire allocated object at allocation-time, which adds linear overhead if the object was not otherwise being initialized.

$$locals\ xs\ m\ le = \begin{cases} (m, le) & \text{if } xs = \varepsilon \\ locals\ xs'\ m''\ le' & \text{if } xs = (id, ty) :: xs' \\ & \text{where } (m', p) \leftarrow stack_alloc\ |ty|\ m, \\ & m'' = m'[p \dots p + |ty| \mapsto \mathbf{undef@vt@}\overline{lt}]_{|ty|}, \\ & pt, vt, \overline{lt} \leftarrow \mathbf{GlobalT}(id, |ty|), \\ & \text{and } le' = le[id \mapsto (p, p + |ty|, ty, pt)] \end{cases}$$

In the tag rule, the PC Tag carries the “next” color to be assigned. We mark both the pointer tag (which is stored in the local environment) with that color, along with the location tags on the allocated memory. Then we increment the PC Tag to give the next allocation a unique color.

LocalT(\mathcal{P}, id, s)

$$\begin{aligned} pt &\leftarrow dyn\ \mathcal{P} \\ \boxed{vt} &\leftarrow X \\ \boxed{\overline{lt}} &\leftarrow [dyn\ \mathcal{P} \mid 0 \leq i < s] \\ \mathcal{P}' &\leftarrow \mathcal{P} + 1 \end{aligned}$$

Heap objects are the most interesting: they are allocated via calls to `malloc`. In Tagged C, `malloc` is modeled as an external call to a built-in, so this takes the form of a special case of that expression. Where *heap_alloc* is some allocation function (a parameter of the memory model) that takes a size and a memory and returns an address:

$$\frac{\mathcal{P}', pt, \boxed{vt}, \boxed{lt} \leftarrow \mathbf{MallocT}(\mathcal{P}, st) \quad m, p \leftarrow \text{heap_alloc } |ty| \ m \quad m'' = m' [p + i \mapsto (\mathbf{undef}, vt, lt) \mid 0 \leq i < s]}{\mathcal{E}(m, ge, le \mid k [\text{builtin malloc}(size@st) : ty] @\mathcal{P}) \longrightarrow \mathcal{E}(m'', ge, le \mid k [p@pt : ty] @\mathcal{P}')$$

And the tag rule is identical to **LocalT**, except that it always treats the allocated object as an array of bytes (making the location tags are always identical.)

$$\begin{aligned} & \mathbf{MallocT}(\mathcal{P}, st) \\ & \mathbf{pt} \leftarrow \text{dyn } \mathcal{P} \\ & \boxed{vt} \leftarrow X \\ & \boxed{\overline{lt}} \leftarrow [\text{dyn } \mathcal{P}] \\ & \mathcal{P}' \leftarrow \mathcal{P} + 1 \end{aligned}$$

Color Checking When we perform a memory load or store, we check that the pointer tag on the left hand of the assignment matches the location tag on all of the bytes being loaded or stored. For instance, in a normal *valof* expression, which accesses a left-hand value:

$$\frac{m[l]_{|ty|} = v@vt@\overline{lt} \quad vt' \leftarrow \mathbf{LoadT}(\mathcal{P}, pt, vt, \overline{lt})}{\mathcal{E}(m, ge, le \mid k [*l@pt : ty] @\mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k [v@vt' : ty] @\mathcal{P})}$$

We want to both check that the pointer tag matches all of the location tags, and propagate the value tag on the value in memory alongside that value.

$$\begin{aligned} & \mathbf{LoadT}(\mathcal{P}, pt, vt, \overline{lt}) \\ & \mathbf{assert} \ \forall lt \in \overline{lt}. pt = lt \\ & vt' \leftarrow vt \end{aligned}$$

There are two other expressions that load from memory, and which therefore invoke this same rule, *assignop* and *postincr*. Note that the C spec has the order of evaluation for *assignop* “unsequenced”; I follow CompCert in evaluating both the left and right completely before performing the load. Intuitively, assignment-with-an-operator is classed along with the standard assignment in the spec, so it is appropriate that it be ordered in the same way.

$$\frac{m[l]_{|ty|} = v_1@vt_1@\overline{lt} \quad \oplus \in \{+, -, *, /, \%, <<, >>, \&, ^, |\} \quad vt'_1 \leftarrow \mathbf{LoadT}(\mathcal{P}, pt, vt_1, \overline{lt})}{\mathcal{E}(m, ge, le \mid k [\underline{l}@pt \oplus] = v_2@vt_2 : ty] @\mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k [\underline{l}@pt := v_1@vt'_1 \oplus v_2@vt_2 : ty] @\mathcal{P})}$$

$$\frac{m[l] = v_1@vt_2@\overline{lt} \quad \oplus \in \{+, -\} \quad vt'_1 \leftarrow \mathbf{LoadT}(\mathcal{P}, pt, vt_1, \overline{lt})}{\mathcal{E}(m, ge, le \mid k [\underline{l}@pt \oplus \oplus : ty] @\mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k [\underline{l}@pt := v@vt'' \oplus 1@\mathbf{ConstT}, v@vt'' : ty] @\mathcal{P})}$$

On the flip side, we store values to memory using the *assign* expression:

$$\frac{m[l]_{|ty|} = v_1@vt_1@\overline{lt} \quad m' = m[l \mapsto v_2@vt'@\overline{lt}'] \quad \mathcal{P}', vt', \overline{lt}' \leftarrow \mathbf{StoreT}(\mathcal{P}, pt, vt_1, vt_2, \overline{lt})}{\mathcal{E}(m, ge, le \mid k [\underline{l}@pt := v_2@vt_2 : ty] @\mathcal{P}) v_2@vt_2 \longrightarrow \mathcal{E}(m', ge, le \mid k [v_2@vt_2 : ty] @\mathcal{P}')$$

As before, we check that the pointer tag matches the locations tags, and then propagate the value tag (ignoring and overwriting the original value tag.) In addition, we propagate the PC Tag.

$$\begin{array}{l}
\mathbf{StoreT}(\mathcal{P}, pt, vt_1, vt_2, \overline{lt}) \\
\text{assert } \forall lt \in \overline{lt}. pt = lt \\
\mathcal{P}' \longleftarrow \mathcal{P} \\
vt' \longleftarrow vt_2 \\
\overline{lt}' \longleftarrow \overline{lt}
\end{array}$$

Color Propagation When a value moves from one location to another, it carries the same tag. We already saw this in the load and store rules: they maintain the relationship between the pointer and its tag. Of note here is the **VarT** control point, which transmits the pointer tag from the environment onto the location expression. In this policy, it propagates the color unchanged.

$$\frac{le[id] = (l, -, pt, ty) \quad pt \leftarrow \mathbf{VarT}(\mathcal{P}, vt)}{\mathcal{E}(m, ge, le \mid k[id : ty] @ \mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k[l @ pt : ty] @ \mathcal{P})}$$

Then the color is propagated via all unary operations and all binary operations where exactly one argument has a color. Performing an operation with two values with color tags (i.e., two cast pointers) clears the tag on the result. It can still be used as an integer, but if cast back to a pointer it will be invalid.

$$\begin{array}{l}
\frac{\langle \odot \rangle v = v' \quad vt' = \mathbf{UnopT}(\mathcal{P}, vt)}{\mathcal{E}(m, ge, le \mid k[\odot v @ vt : ty] @ \mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k[v' @ vt' : ty] @ \mathcal{P})} \\
\frac{v_1 \langle \oplus \rangle v_2 = v' \quad vt' = \mathbf{BinopT}(\mathcal{P}, vt_1, vt_2)}{\mathcal{E}(m, ge, le \mid k[v_1 @ vt_1 \oplus v_2 @ vt_2 : ty] @ \mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k[v' @ vt' : ty] @ \mathcal{P})}
\end{array}$$

$$\begin{array}{ll}
\mathbf{UnopT}(\mathcal{P}, vt) & \mathbf{BinopT}(\mathcal{P}, vt_1, vt_2) \\
\mathcal{P}' \longleftarrow \mathcal{P} & \mathcal{P}' \longleftarrow \mathcal{P} \\
vt' \longleftarrow vt & vt' \longleftarrow \text{case } (vt_1, vt_2) \text{ of} \\
& \quad \text{dyn } n, X \Rightarrow \text{dyn } n \\
& \quad \text{glob } id, X \Rightarrow \text{glob } id \\
& \quad X, t \Rightarrow t
\end{array}$$

3.2 PNVI Definitions

In PNVI, the basic provenance model remains the same as PVI, so we can reuse most of the same rules. The primary difference is what happens when we cast a pointer to an integer. In PVI, tags are propagated as normal.

To support PNVI, we need the *cast* expression to update the tags of a pointer being cast to an integer and vice versa. We add two special-case steps to reflect this.

$$\frac{\boxed{m[p]_{|ty|} = _ @ vt_2 @ \overline{lt}} \quad \mathcal{P}', vt \leftarrow \mathbf{PICastT}(\mathcal{P}, pt, \boxed{vt, \overline{lt}})}{\mathcal{E}(m, ge, le \mid k[(int)p @ pt : ptr(ty)] @ \mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k[p @ vt : int] @ \mathcal{P})}$$

$$\frac{\boxed{m[p]_{|ty|} = _@vt_2@\overline{lt}} \quad \mathcal{P}', pt \leftarrow \mathbf{IPCastT}(\mathcal{P}, vt_1, \boxed{vt_2, \overline{lt}})}{\mathcal{E}(m, ge, le \mid k[(int)p@pt : ptr(ty)] @ \mathcal{P}) \longrightarrow \mathcal{E}(m, ge, le \mid k[p@vt : int] @ \mathcal{P})}$$

For casting an integer to a pointer, we don't need the optional "peek" at the memory that it points to. We simply clear the tag on the resulting integer.

$$\begin{aligned} & \mathbf{PICastT}(\mathcal{P}, pt, \boxed{-, -}) \\ \mathcal{P}' & \longleftarrow \mathcal{P} \\ vt & \longleftarrow X \end{aligned}$$

On the other hand, when casting back to a pointer, we need to check the color of the object that it points to.

$$\begin{aligned} & \mathbf{IPCastT}(\mathcal{P}, vt_1, \boxed{vt_2, \overline{lt}}) \\ & \quad \mathbf{assert} \ \exists t. \forall lt \in \overline{lt}. lt = t \wedge t \neq X \\ \mathcal{P}' & \longleftarrow \mathcal{P} \\ pt & \longleftarrow t \end{aligned}$$

Realizing the Integer-Pointer Cast The pointer cast rules take as input the tags on the location pointed to by the argument being cast. This requires the compiler to add extra instructions to retrieve that tag. On RISC-V, the sequence would be as follows, assuming that `a0` contains the value being cast. The meaning of instruction tags will be explained below.

```
lw a1 a0 0 @ RETRIEVE
sub a1 a1 a1 @ L
add a0 a1 a0 @ IPCAST
```

In the underlying assembly, we use instruction tags to inform the low-level monitor of the purpose of each instruction. **RETRIEVE** indicates a special load whose job is retrieve value and location tags from a location in memory. When it sees a **RETRIEVE** tag, the monitor allows the load even if it should failstop under the Concrete C backstop policy. If the load should failstop, however, it is given a default tag rather than the tags on the memory. A legal load receives both the value and the location tags.

The **L** instruction tag simply denotes taking the left-operand's tag on the result of a binary operation. In this case both operations are identical, but we still need to pick one. Finally, the **IPCAST** tag declares that this instruction should mimic the Tagged-C-level rule.

4 Information Flow Control

Next we will discuss *information flow control* (IFC), specifically *partial IFC* (pIFC) or *taint-tracking*. In pIFC, we define a limited set of rules for how information is allowed to flow from *sources* to *sinks*. A source σ can be an argument of a function or its return value. A sink ψ can be a global variable, a function, or the set of heap objects allocated by a given function. If we identify a source as carrying high-privilege secrets, and a sink as lacking the privilege to receive those secrets, then we would define a *no-flow* rule between them, written $\sigma \not\rightsquigarrow \psi$. Likewise if the source is untrusted, and

the sink highly sensitive. Additionally, an IFC policy may contain “declassification” rules, σ/σ' , which means that if a value comes from the source σ' , we may safely ignore its original source σ . We will often write $*/\sigma$ to say that σ declassifies anything. We model multiple sources and sinks simultaneously, and absent a no-flow rule, they can interact freely.

For example, suppose that in the following code, we have a no-flow rule between the argument x of f and the global variable z ($f.x \not\rightarrow z$), and a declassification rule $*/g.a$.

```
int z;

int g(int a);

void f(int x, int y) {
    z = x;           // violation
    z = x + y;       // violation
    if(x) z = 1; else z = 0; // violation
    z = g(x);        // legal
}
```

The first three lines of f violate the no-flow relation by storing values derived from x into z . The third line is especially interesting: although x is not stored directly, the value that is stored is conditioned upon it, and can be used to deduce information about the original value. This is termed an *implicit flow*. Finally, in the last line, while the value of $g(x)$ may depend on x , it is subject to the declassification rule, and so it is permissible.

Tagging Taint We track the influence of a particular source, or “taint,” through the system in the form of tags on values. We identify sinks through static tags on functions and through memory location tags, that encode the “forbidden” sources. So, our tag set is as follows, with X being the default tag that indicates no (interesting) source has tainted the value. Note that our tags may carry multiple sources.

$$\begin{aligned} T ::= & \text{taint } \bar{\sigma} \\ & \text{forbid } \bar{\sigma} \\ & X \end{aligned}$$

We define three important operations on tags: *merge* (\sqcup), *minus* ($-$), and *check* (\sqsubseteq), all partial functions.

$$\begin{aligned} t_1 \sqcup t_2 &\triangleq \begin{cases} \text{taint } (\bar{\sigma}_1 \cup \bar{\sigma}_2) & \text{if } t_1 = \text{taint } \bar{\sigma}_1 \text{ and } t_2 = \text{taint } \bar{\sigma}_2 \\ \perp & \text{otherwise} \end{cases} \\ t - \sigma &\triangleq \begin{cases} \text{taint } (\bar{\sigma}' - \sigma) & \text{if } t = \text{taint } \bar{\sigma}' \\ \perp & \text{otherwise} \end{cases} \\ t_1 \sqsubseteq t_2 &\triangleq \begin{cases} \mathbf{t} & \text{if } t_1 = \text{taint } \bar{\sigma}_1, t_2 = \text{forbid } \bar{\sigma}_2, \text{ and } \bar{\sigma}_1 \cap \bar{\sigma}_2 = \emptyset \\ \mathbf{f} & \text{else if } t_1 = \text{taint } \bar{\sigma}_1 \text{ and } t_2 = \text{forbid } \bar{\sigma}_2 \\ \perp & \text{otherwise} \end{cases} \end{aligned}$$

Introducing Taint Each no-flow rule specifies a source that is either a function argument or return value. These attach to the *call-state* and *return-state* rules, respectively. The call-state rule executes at the beginning of a call, moving all of its arguments into the local environment, using the **ArgT** tag rule.

$$\frac{\text{def}(f) = (xs, s) \quad \text{le}' = \text{le} \llbracket x \mapsto v @ vt' \mid (x, v @ vt) \leftarrow \text{zip}(xs, args), vt' \leftarrow \mathbf{ArgT}_I(\mathcal{P}, vt, x) \rrbracket}{\mathcal{C}(m, ge, le \mid k[f(args)] @ \mathcal{P}) \longrightarrow \mathcal{S}(m, ge, le' \mid s;; k @ \mathcal{P})}$$

And the return-state rule executes after the call returns, inserting the result into the context saved in the continuation. The program counter on return and the result's tag are set by the **CallerRetT** tag rule.

$$\frac{k = Kcall \text{ le}' \text{ ctx } k' \quad \mathcal{P}'', vt' \leftarrow \mathbf{CallerRetT}(\mathcal{P}, \mathcal{P}', vt)}{\mathcal{R}(m, ge, le \mid k[v @ vt] @ \mathcal{P}) \longrightarrow \mathcal{S}(m, ge, le' \mid \text{ctx}[v @ vt'];; k' @ \mathcal{P}')$$

In the case of an IFC policy, both control points are parameterized by a set of IFC rules, I .

$$\begin{aligned} & \mathbf{ArgT}_I(\mathcal{P}, vt, nt) \\ & \text{let } vt' := vt - \{\sigma \mid \sigma/nt \in I\} \text{ in} \\ & \text{let } vt'' := vt' \sqcup \text{tainted } \{nt \mid nt \not\prec _ \in I\} \text{ in} \\ & vt' \longleftarrow vt'' \end{aligned}$$

Propagating Taint Through Expressions It is simple enough to determine when a value is tainted: at a function call, all function arguments are tagged with their source identity, and the result of any expression is tagged with the union of the sources of its operands. If the expression involves a store or function call itself, we must check the taints on the value being stored or passed against the forbidden list of the target.

Unary and binary operations:

$$\begin{array}{ll} \mathbf{UnopT}(\mathcal{P}, vt) & \mathbf{BinopT}(\mathcal{P}, vt_1, vt_2) \\ \mathcal{P}' \longleftarrow \mathcal{P} & \mathcal{P}' \longleftarrow \mathcal{P} \\ vt' \longleftarrow vt & vt' \longleftarrow vt_1 \sqcup vt_2 \end{array}$$

Loads and stores:

$$\begin{array}{ll} \mathbf{LoadT}(\mathcal{P}, pt, vt, \overline{lt}) & \mathbf{StoreT}(\mathcal{P}, pt, vt_1, vt_2, \overline{lt}) \\ vt' \longleftarrow \mathcal{P} \sqcup pt \sqcup vt & \mathbf{assert} \forall lt \in \overline{lt}. \mathcal{P} \sqcup pt \sqcup vt_2 \sqsubseteq lt \\ & \mathcal{P}' \longleftarrow \mathcal{P} \\ & vt' \longleftarrow \mathcal{P} \sqcup pt \sqcup vt_2 \\ & \overline{lt}' \longleftarrow \overline{lt} \end{array}$$

Implicit Flows Things become trickier when control-flow itself can be tainted. This can occur in any of our semantics' steps that can produce different statements and continuations depending on the tainted value. At that point, any change to the machine state constitutes an information flow.

```

int f(bool secret) {
    int public1, public2;

S:  if (secret) {
b1:    public1 = 1;
    } else {
b2:    public1 = 0;
    }

J:   public2 = 42;

    return public2;
}

```




Figure 2: Leaking via if statements

To be more specific, consider a statement that contains an expression, $s(e)$, such that when filled in with a tainted value:

$$\mathcal{S}(m, ge, le \mid sv_1 @ \text{taint } \sigma;; k @ \mathcal{P}) \longrightarrow \mathcal{S}(m_1, ge_1, le_1 \mid s_1;; k_1 @ \mathcal{P}_1)$$

while

$$\mathcal{S}(m, ge, le \mid sv_1 @ \text{taint } \sigma;; k @ \mathcal{P}) \longrightarrow \mathcal{S}(m_2, ge_2, le_2 \mid s_2;; k_2 @ \mathcal{P}_2)$$

and where $s_1 \neq s_2$ or $k_1 \neq k_2$. Taking either step should taint the program state itself! We represent this as a taint on the PC Tag. When the PC Tag is tainted, all stores to memory and all updates to environments must also be tainted.

This presents a problem if the program counter must remain tainted indefinitely. Fortunately, it is safe to remove the taint if all branches eventually rejoin. We term this a *join point*. In terms of the program's control-flow graph, the join point of a branch is its immediate post-dominator.

In many simple programs, the join point of a conditional or loop is obvious: the point at which the chosen branch is complete, or the loop has ended. Such a simple example can be seen in fig. 2; `public1` must be tagged with the taint tag of `secret`, while it is safe to tag `public2` *X*, because that is after the join point, J. The same goes for fig. 3.

But in the presence of unrestricted go-to statements, a join point may not be local—and sometimes may not exist. Consider fig. 4, which uses go-to statements to create an approximation of an if-statement whose join-point is far removed from the for-loop. The label J now has nothing to do with the semantics of any particular statement.

Luckily this can still be determined statically from a function's full control-flow graph. So, we annotate our programs with this information via additional labels. Every statement that branches can carry an optional label indicating its corresponding join point. If it doesn't have such a label, that indicates that there is no join point—once the PC Tag is tainted, it must remain so until a return.

When we step into a conditional or loop, we need to record its join point for later. The join point will always be represented by a label, which is an identifier, so we can simply store it in the local environment.

```

int f(bool secret) {
    int public1=1;
    int public2;

    S: while (secret) {
        b1: public1 = 1;
            secret = false;
        }

    J: public2 = 42;

        return public2;
    }

```




Figure 3: Leaking via while statements

```

int f(bool secret) {
    int public1, public2;

    while (secret) {
        goto b1;
    }

    b2: public1 = 1;
        goto J;

    b1: public1 = 1;

    J: public2 = 42;
        return public2;
    }

```




Figure 4: Cheating with go-tos

$$\frac{s' = \begin{cases} s_1 & \text{if } \text{boolof}(v) = \mathbf{t} \\ s_2 & \text{if } \text{boolof}(v) = \mathbf{f} \end{cases} \quad \begin{array}{l} \mathcal{P}'', \mathcal{P}''' \leftarrow \mathbf{SplitT}(\mathcal{P}, vt, \mathcal{P}') \\ le[L] = _@ \mathcal{P}' \quad le' = le[L \mapsto \mathbf{undef}@ \mathcal{P}'''] \end{array}}{\mathcal{S}(m, ge, le \mid \text{if}(v@vt) \text{ then } s_1 \text{ else } s_2 \text{ join } L;; k@ \mathcal{P}) \longrightarrow \mathcal{S}(m, ge, le' \mid s';; k@ \mathcal{P}')}$$

The label carries a record of all sources that might currently taint the program's control flow, but will be safe once that label is reached. So, just as \mathcal{P} is merged with vt to produce \mathcal{P}'' , \mathcal{P}' will be merged with vt to produce \mathcal{P}''' , a new tag for L . Then, when we reach the label L , we will retrieve the stored tag and subtract its sources from the PC Tag at that time.

$$\begin{array}{ll} \mathbf{SplitT}(\mathcal{P}, vt, \boxed{\mathcal{P}'}) & \mathbf{JoinT}(\mathcal{P}, \boxed{\mathcal{P}'}) \\ \mathcal{P}'' \leftarrow \mathcal{P} \sqcup vt & \text{let } \mathcal{P}'' := \mathcal{P} - \{\sigma \mid \sigma'\} \text{ in} \\ \boxed{\mathcal{P}'''} \leftarrow \mathcal{P}' \sqcup vt & \mathcal{P}'' \leftarrow \mathcal{P} \sqcup vt \\ & \boxed{\mathcal{P}'''} \leftarrow \mathcal{P}' \sqcup vt \end{array}$$

The **JoinT** control point applies whenever we reach a labeled statement, like so:

$$\frac{le[L] = _@ \mathcal{P}' \quad \mathcal{P}'', \mathcal{P}''' \leftarrow \mathbf{JoinT}(\mathcal{P}, \mathcal{P}') \quad le' = le[L \mapsto \mathbf{undef}@ \mathcal{P}''']}{\mathcal{S}(m, ge, le \mid L : s;; k@ \mathcal{P}) \longrightarrow \mathcal{S}(m, ge, le' \mid s;; k@ \mathcal{P}'')}$$

The remaining branching constructs are rather complicated, involving multiple steps and manipulations of the continuation that are not that relevant to their control points. Rather than give their semantics in full, it suffices to identify which transitions contain **SplitT** control points. In fig. 5, these are the transitions from the state marked S .

Realizing IFC In order to implement an IFC policy, we need to specify the rules that it needs to enforce. The positive here is that the rules are not dependent on one another (with the exception of declassification rules), and default to permissiveness when no rule is given. We assume that the user would supply a separate file consisting of a list of triples: the source, the sink, and the type of rule. This is then translated into the policy.

The other implementation detail to consider are the label tags. These resemble instruction tags, and that is exactly how they would be implemented: as a special instruction tag on the appropriate instruction, which might be an existing instruction or a specially added no-op. But importantly, in this case, these tags are mutable; in a policy that can be expected to take advantage of their mutability, we will need an extra store to set the tag for later.

It remains to generate those labels. For purposes of an IFC policy, we first generate the program's control flow graph. Then, for each if, while, do-while, for, and switch statement, we identify the immediate post-dominator in the graph, and wrap it in a label statement with a fresh identifier. That identifier is also added as a field in the original conditional statement. The tags associated with the labels are initialized at program state—in the case of IFC, these defaults declare that there are no secrets to lowre when it is reached.

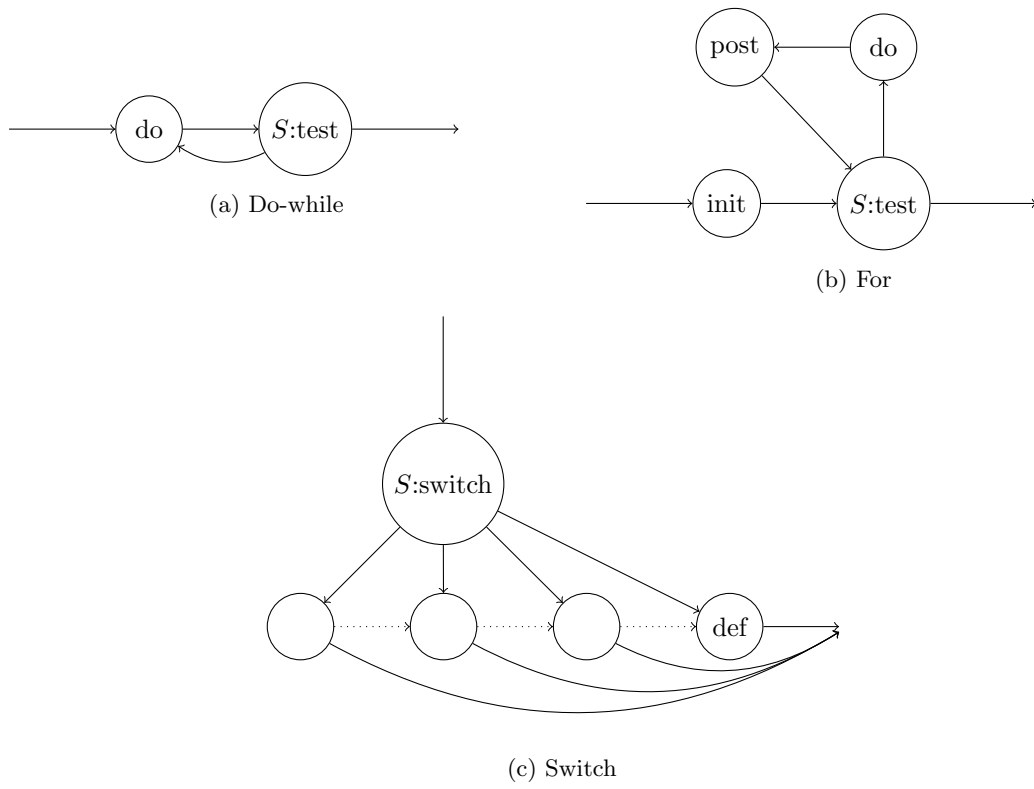


Figure 5: Remaining Branch Statements